



# **Agilio OvS 2.2 Programmer's Reference Manual**

## **Netronome Intelligent Server Adapters**

**- Proprietary and Confidential -**

# Agilio OvS 2.2 Programmer's Reference Manual: Netronome Intelligent Server Adapters

Copyright © 2011-2016 Netronome

## COPYRIGHT

No part of this publication or documentation accompanying this Product may be reproduced in any form or by any means or used to make any derivative work by any means including but not limited to by translation, transformation or adaptation without permission from Netronome Systems, Inc., as stipulated by the United States Copyright Act of 1976. Contents are subject to change without prior notice.

## WARRANTY

Netronome warrants that any media on which this documentation is provided will be free from defects in materials and workmanship under normal use for a period of ninety (90) days from the date of shipment. If a defect in any such media should occur during this 90-day period, the media may be returned to Netronome for a replacement.

NETRONOME DOES NOT WARRANT THAT THE DOCUMENTATION SHALL BE ERROR-FREE. THIS LIMITED WARRANTY SHALL NOT APPLY IF THE DOCUMENTATION OR MEDIA HAS BEEN (I) ALTERED OR MODIFIED; (II) SUBJECTED TO NEGLIGENCE, COMPUTER OR ELECTRICAL MALFUNCTION; OR (III) USED, ADJUSTED, OR INSTALLED OTHER THAN IN ACCORDANCE WITH INSTRUCTIONS FURNISHED BY NETRONOME OR IN AN ENVIRONMENT OTHER THAN THAT INTENDED OR RECOMMENDED BY NETRONOME.

EXCEPT FOR WARRANTIES SPECIFICALLY STATED IN THIS SECTION, NETRONOME HEREBY DISCLAIMS ALL EXPRESS OR IMPLIED WARRANTIES OF ANY KIND, INCLUDING, WITHOUT LIMITATION, ANY IMPLIED WARRANTIES OF MERCHANTABILITY, NON-INFRINGEMENT AND FITNESS FOR A PARTICULAR PURPOSE.

Some jurisdictions do not allow the exclusion of implied warranties, so the above exclusion may not apply to some users of this documentation. This limited warranty gives users of this documentation specific legal rights, and users of this documentation may also have other rights which vary from jurisdiction to jurisdiction.

## LIABILITY

Regardless of the form of any claim or action, Netronome's total liability to any user of this documentation for all occurrences combined, for claims, costs, damages or liability based on any cause whatsoever and arising from or in connection with this documentation shall not exceed the purchase price (without interest) paid by such user.

IN NO EVENT SHALL NETRONOME OR ANYONE ELSE WHO HAS BEEN INVOLVED IN THE CREATION, PRODUCTION, OR DELIVERY OF THE DOCUMENTATION BE LIABLE FOR ANY LOSS OF DATA, LOSS OF PROFITS OR LOSS OF USE OF THE DOCUMENTATION OR FOR ANY SPECIAL, INCIDENTAL, CONSEQUENTIAL, EXEMPLARY, PUNITIVE, MULTIPLE OR OTHER DAMAGES, ARISING FROM OR IN CONNECTION WITH THE DOCUMENTATION EVEN IF NETRONOME HAS BEEN MADE AWARE OF THE POSSIBILITY OF SUCH DAMAGES. IN NO EVENT SHALL NETRONOME OR ANYONE ELSE WHO HAS BEEN INVOLVED IN THE CREATION, PRODUCTION, OR DELIVERY OF THE DOCUMENTATION BE LIABLE TO ANYONE FOR ANY CLAIMS, COSTS, DAMAGES OR LIABILITIES CAUSED BY IMPROPER USE OF THE DOCUMENTATION OR USE WHERE ANY PARTY HAS SUBSTITUTED PROCEDURES NOT SPECIFIED BY NETRONOME.

## Revision History

Date	Revision	Description
24 February 2016	000	First DRAFT for Agilio OvS 2.2.

# Table of Contents

<b>1. Introduction</b>	7
1.1. Scope	7
1.2. Related Documents	7
1.3. Terminology	7
<b>2. Local Flow API</b>	9
2.1. Purpose	9
2.2. Overview	9
2.3. Table Types	10
2.4. Sample Applications	11
2.4.1. CSV Sample Application	11
2.4.2. Command Line Control Sample Application	11
<b>3. Group API</b>	12
3.1. Purpose	12
3.2. Overview	12
3.3. Sample Application	12
<b>4. Health Monitoring API</b>	13
4.1. Purpose	13
4.2. Overview	13
4.3. Sample Applications	13
<b>5. Agilio OVS 2.2 Host API Reference</b>	14
5.1. ofp_flow	14
5.1.1. Defines	14
5.1.2. Enumerations	14
5.1.3. Typedefs	15
5.1.4. Functions	16
5.2. ofp_group	30
5.2.1. Enumerations	30
5.2.2. Functions	31
5.3. ofp_action	33
5.3.1. Defines	33
5.3.2. Enumerations	33
5.3.3. Functions	34
5.4. ns_sdn_health	41
5.4.1. Defines	41
5.4.2. Enumerations	42
5.4.3. Typedefs	43
5.4.4. Functions	43
<b>6. Technical Support</b>	48
6.1. Related Documents	48

## List of Tables

5.1. ofp_flow.h Defines .....	14
5.2. enum ofp_flow_mod_command .....	14
5.3. enum ofp_flow_table_type .....	14
5.4. enum ofp_group_no .....	15
5.5. enum ofp_port_no .....	15
5.6. enum ofp_table .....	15
5.7. typedef openflow_error_fptr_t .....	16
5.8. ofp_sdn_sys_init parameters .....	16
5.9. ofp_sdn_barrier parameters .....	17
5.10. ofp_sdn_barrier_timeout parameters .....	18
5.11. ofp_flow_entry_create parameters .....	19
5.12. ofp_flow_entry_free parameters .....	20
5.13. ofp_flow_entry_match_field parameters .....	20
5.14. ofp_flow_entry_match_field_mask parameters .....	21
5.15. ofp_flow_entry_cookie parameters .....	21
5.16. ofp_flow_entry_xid parameters .....	22
5.17. ofp_flow_entry_hard_timeout parameters .....	22
5.18. ofp_flow_entry_idle_timeout parameters .....	23
5.19. ofp_flow_entry_priority parameters .....	23
5.20. ofp_flow_entry_out_port parameters .....	24
5.21. ofp_flow_entry_out_group parameters .....	25
5.22. ofp_flow_entry_table_num parameters .....	25
5.23. ofp_flow_entry_instr_goto_table parameters .....	26
5.24. ofp_flow_entry_instr_write_metadata parameters .....	26
5.25. ofp_flow_entry_instr_clear_actions parameters .....	27
5.26. ofp_flow_entry_instr_meter parameters .....	27
5.27. ofp_flow_entry_instr_write_actions parameters .....	28
5.28. ofp_flow_entry_instr_apply_actions parameters .....	29
5.29. ofp_flow_entry_commit parameters .....	29
5.30. ofp_flow_entry_stats parameters .....	30
5.31. enum ofp_group_mod_command .....	30
5.32. enum ofp_group_type .....	30
5.33. ofp_group_create parameters .....	31
5.34. ofp_group_add_action_bucket parameters .....	31
5.35. ofp_group_commit parameters .....	32
5.36. ofp_group_free parameters .....	32
5.37. ofp_action.h Defines .....	33
5.38. enum ofp_action_mode .....	33
5.39. enum oxm_ofb_match_fields .....	33
5.40. ofp_action_create parameters .....	35
5.41. ofp_action_free parameters .....	35
5.42. ofp_action_output parameters .....	36
5.43. ofp_action_group parameters .....	36
5.44. ofp_action_mpls_ttl parameters .....	37
5.45. ofp_action_dec_mpls_ttl parameters .....	37
5.46. ofp_action_ip_ttl parameters .....	37
5.47. ofp_action_dec_ip_ttl parameters .....	38
5.48. ofp_action_push_mpls parameters .....	38
5.49. ofp_action_pop_mpls parameters .....	39
5.50. ofp_action_set_queue parameters .....	39
5.51. ofp_action_push_vlan parameters .....	40

5.52. ofp_action_pop_vlan parameters .....	40
5.53. ofp_action_set_field parameters .....	41
5.54. ns_sdn_health.h Defines .....	41
5.55. enum ns_sdn_health_failure_code_t .....	42
5.56. enum ns_sdn_health_optional_checks_t .....	43
5.57. typedef ns_sdn_health_failure_t .....	43
5.58. ns_sdn_health_init parameters .....	43
5.59. ns_sdn_health_destroy parameters .....	44
5.60. ns_sdn_health_check parameters .....	44
5.61. ns_sdn_health_failure_count parameters .....	45
5.62. ns_sdn_health_warning_count parameters .....	45
5.63. ns_sdn_health_next_failure parameters .....	46
5.64. ns_sdn_health_generate_report parameters .....	46

# 1. Introduction

---

## 1.1 Scope

This manual details the host APIs for programmatically manipulating aspects of the Netronome Systems Agilio OvS 2.2 product. The intended audience for this book includes developers and system programmers.

This manual is organized as follows:

- Chapter 2 provides a high-level description of the local Flow API for manipulating the Agilio OvS 2.2 Flow Tables.
- Chapter 3 provides a high-level description of the group API for manipulating the Agilio OvS 2.2 group entries.
- Chapter 4 provides a high-level description of the health API.
- Chapter 5 provides a complete reference of the APIs.
- Chapter 6 provides information on technical support.

## 1.2 Related Documents

Descriptive Name	Description
Netronome Agilio OvS 2.2 User's Guide	Provides a general overview of the Agilio OvS firmware.
Netronome Agilio OvS 2.2 Getting Started Guide	A guide to new users of Netronome's OpenFlow switch software.
OpenFlow Switch Specification 1.3.3	Control protocol standardized by Open Networking Foundation.

## 1.3 Terminology

Acronym	Description
API	Application Programming Interface
CIDR	Classless Inter-Domain Routing
CSV	Comma Separated Values, a simple text file format
FST	Flow State Table
IPv4	Internet Protocol version 4
LB	Load Balancing

Acronym	Description
MAC	Media Access Control
MPLS	MultiProtocol Label Switching
NFD	Network Flow Driver
NFP	Netronome Network Flow Processor
NIC	Network Interface Card
OvS	Open vSwitch
PCIe	PCI Express
SDN	Software Defined Networking
TCP	Transmission Control Protocol
TTL	Time To Live
VLAN	Virtual Local Area Network
VM	Virtual Machine
VNIC	Virtual Network Interface Card



## 2. Local Flow API

### 2.1 Purpose

The local Flow API allows the user to manipulate (add, modify, delete) flow entries in Netronome Agilio OvS 2.2 flow tables. These flow entry operations closely mirror OpenFlow semantics as described in the OpenFlow 1.3 specification<sup>1</sup>.

### 2.2 Overview

A flow entry typically consists of the following information:

Match Fields	Priority	Counters	Instructions	Timeouts	Cookie	Flags
--------------	----------	----------	--------------	----------	--------	-------

The API function prototypes are declared in the `ofp_flow.h` file<sup>2</sup>, and these functions allow the user to add, modify and delete flow entries in specified tables. The flow entry fields (as in the table definition above, with the exception of Counters) can also be manipulated by these APIs. During flow processing, the Match Fields are compared against packet headers, ingress port and other fields such as metadata from a previous table match. The Priority specifies the precedence of the flow entry. A flow table entry is uniquely identified by the Match Fields and Priority combination. The flow entry that wildcards all fields (all fields omitted) and has priority set to 0 is called the table-miss flow entry. The Cookie field is an opaque value chosen by the user, not used when processing packets but used to later filter flow entries for statistics, flow modification and flow removal.

The Instructions field consists of an ordered set comprising at most one of each of the following:

- Meter *meter\_id*
- Apply-Actions *action\_list*
- Clear-Actions
- Write-Actions *action\_set*
- Write-Metadata *metadata/mask*
- Goto-Table *next-table-id*

APIs are provided to enable each of these instructions in the flow entry. The APIs for the actions referred to by the two action instructions need to be called in an apply or write mode, to distinguish between the ordered list required by the Apply-Actions instruction, and the set required by the Write-Actions instruction. To this end, the user must call `ofp_action_create()` prior to calling a series of `ofp_action_*` functions, and finally the apply or write instruction can be enabled. The general usage of the API is as follows:

1. Initialize an `ofp_flow` subsystem handle ( `ofp_sdn_sys_init()` ). An error callback function pointer may be registered, this will be invoked in the event of an error indication from the switch, and provides a transaction id, error type and error code (details can found in the OpenFlow specification).

<sup>1</sup>Refer to the section named "OpenFlow Tables" in the OpenFlow Switch Specification 1.3.3 on the [www.opennetworking.org](http://www.opennetworking.org) website.

<sup>2</sup>This file is included in the Agilio OvS 2.2 installation, it can be found in the `/opt/netronome/include` directory.

2. Create an `ofp_flow_entry` handle ( `ofp_flow_entry_create()` ), specifying table type, whether the entry is for add, delete or modify operation, and operation-specific flags. Some table types do not support the full OpenFlow semantics (see Section 2.3).
3. Set the desired packet match fields. If left unspecified, the flow entry will match all packets ( `ofp_flow_entry_match_*` APIs).
4. Set the priority, defaults to 0 if left unspecified ( `ofp_flow_entry_priority()` ). Note: Omitting both steps 3 and 4 will result in the flow entry being regarded as the table-miss flow entry. This special entry specifies how to process packets unmatched by other flow entries in the table.
5. Set the timeouts and cookie fields as desired ( `ofp_flow_entry_timeout()` , `ofp_flow_entry_cookie()` ).
6. Add actions, after creating a new handle ( `ofp_action_create()` ), to be used next with Write-Actions or Apply-Actions instruction ( `ofp_action_*` APIs).
7. Add instructions to the entry ( `ofp_flow_entry_instr_*` APIs).
8. Optionally, a transaction ID may be specified ( `ofp_flow_entry_xid()` ). This is an arbitrary user defined number which can be meaningful in the event of an error callback on a particular flow entry operation. If left unspecified it will auto-increment from the previous value (starting at 0 for the very first transaction).
9. Commit the entry to a specified table number ( `ofp_flow_entry_commit()` ).
10. Repeat steps 2-8 as often as necessary.
11. Terminate the `ofp_flow` subsystem handle ( `ofp_sdn_sys_shutdown()` ).

## 2.3 Table Types

Flow entries are created in numbered tables of different types in Agilio OvS 2.2 (see `ofp_flow_entry_create()` ). Two table types are enumerated:

1. Open vSwitch table via OpenFlow protocol.
2. NFP hardware flow cache table (not implemented currently).

The Open vSwitch tables are programmed via the OpenFlow protocol over a TCP connection. The API library acts as an OpenFlow controller, and Open vSwitch needs to be configured to allow so-called passive TCP connections:

```
ovs-vsctl set-controller br0 tcp:6653
```

In particular, the API uses OpenFlow version 1.3, and Open vSwitch needs to be configured to not exclude that version. This can be accomplished by not restricting the protocol versions at all (which is also the default setting):

```
ovs-vsctl set Bridge br0 protocols=[]
```

or specifically setting it to allow version 1.3:

```
ovs-vsctl set Bridge br0 protocols=OpenFlow13
```

## 2.4 Sample Applications

Sample applications using the Flow API are included in the Agilio OvS 2.2 installation. The source files are located in `/opt/netronome/samples/user_api` and can be compiled using the following command:

```
make
```

### 2.4.1 CSV Sample Application

The application (source in `ofp_flow_csv.c`) adds OvS table 0 entries from a fixed format CSV file containing entries for matching source IPv4 address, destination IPv4 address and an output port action (referring to OpenFlow port number). The IPv4 addresses may optionally include netmasks in CIDR notation or dotted-decimal format, e.g. `192.168.0.0/16` or `192.168.0.0/255.255.0.0` would equivalently match only the top 16 bits of the IP address. The output port (3rd item in each CSV entry) can be set to 0 to indicate drop (i.e. no output action), or -1 to indicate switch normal processing (in effect, L2 learning with Open vSwitch).

The default table-miss rule can be specified by adding a CSV entry containing `0.0.0.0/0` for source and destination IP address, with output port set as desired (or 0 to drop).

### 2.4.2 Command Line Control Sample Application

This sample (`ofp_flow_ctl.c`) provides more exhaustive coverage of the Flow API by implementing a set of command line options to manipulate (add, delete or modify) the flow tables. The complete usage information can be obtained by typing: `ofp_flow_ctl -?`

The user can for example add the table-miss flow entry (by omitting match fields and priority) to output to the OpenFlow NORMAL port as follows:

```
ofp_flow_ctl --add --write-actions output:0xfffffffffa
```

or add a flow entry with masked source/destination IPv4 address (the field 'ip' expands internally to match ethernet type 0x0800):

```
ofp_flow_ctl --add --match-fields ip,ip_src=10.0.0.0/8,ip_dst=10.0.0.0/8 \
--write-actions output:0xfffffffffa
```

A slightly more complete example might set up a chain of two table entries, the first entry (with a timeout) matching some IP range and setting metadata, and the second matching on that metadata and performing an output action to OpenFlow port 1:

```
ofp_flow_ctl --add --table_num 0 --match-fields ip,ip_src=10.0.0.0/8,ip_dst=10.0.0.0/8 \
--priority 5 --timeout 60 --cookie 1234 --write-metadata 42 --goto-table 1
ofp_flow_ctl --add --table_num 1 --match-fields metadata=42 --write-actions output:1
```

## 3. Group API

---

### 3.1 Purpose

The Group API allows the user to manipulate (add, modify, delete) group entries in Netronome Agilio OvS 2.2 group tables. These operations closely mirror OpenFlow semantics as described in the OpenFlow 1.3 specification<sup>1</sup>.

### 3.2 Overview

A group entry typically consists of the following information:

Identifier	Type	Counters	Action buckets
------------	------	----------	----------------

The API function prototypes are declared in the `ns_group.h` file<sup>2</sup>, and these functions allow the user to add, modify and delete group entries.

The general usage of the API is as follows:

1. Create an `ns_group` handle ( `ofp_group_create()` ), specifying the command (add, modify or delete), group type (all, select, indirect, fast failover), and group identifier.
2. Create a new action set ( `ofp_action_create()` ), populate it with desired actions.
3. Add the action set to the group as an action bucket ( `ofp_group_add_action_bucket()` ) with a specified weight.
4. Repeat steps 2-3 as often as necessary. Indirect groups may only have a single action bucket.
5. Commit the group modifier ( `ofp_group_commit()` ).

### 3.3 Sample Application

A minimal Group API sample application is included in the Agilio OvS 2.2 installation. The source file is located in `/opt/netronome/samples/user_api/ofp_group_tst.c` and can be compiled using the following command:

```
make
```

The sample creates a load balancing group outputting to OpenFlow ports 2 and 3.

---

<sup>1</sup>Refer to the section named "OpenFlow Tables" in the OpenFlow Switch Specification 1.3.3 on the [www.opennetworking.org](http://www.opennetworking.org) website.

<sup>2</sup>This file is included in the Agilio OvS 2.2 installation, it can be found in the `/opt/netronome/include` directory.

## 4. Health Monitoring API

---

### 4.1 Purpose

The health monitoring API allows the user to monitor the system health on the host operating system. The health monitoring API allows the user to create their own health monitoring application suitable to the user's operating environment.

### 4.2 Overview

The API function prototypes are declared in the `ns_sdn_health.h` file<sup>1</sup>, and these functions allow the user to create an application to monitor the health of the Agilio OvS system.

The general usage of the API is as follows:

1. Initialize an `ns_sdn_health_handle_t` instance.
2. Execute (`ns_sdn_health_check()`) to run all the health checks on the system. This function can be called periodically if needed.
3. After the health checks have been executed, the results can be processed by calling (`ns_sdn_health_next_failure()`) successively until no more failures are returned.
4. Destroy the `ns_sdn_health_handle_t` instance (`ns_sdn_health_destroy()`).

### 4.3 Sample Applications

A sample applications using the Health Monitoring API are included in the Agilio OvS 2.2 installation. The source files are located in `/opt/netronome/samples/sdn_health` and can be compiled using the following command:

```
make
```

The included sample provide a few possible options to display the current system health. Refer to the usage output of the sample.

---

<sup>1</sup>This file is included in the Agilio OvS 2.2 installation, it can be found in the `/opt/netronome/include` directory.

## 5. Agilio OVS 2.2 Host API Reference

### 5.1 ofp\_flow

#### 5.1.1 Defines

Table 5.1. ofp\_flow.h Defines

Defined	Definition
ofp_flow_entry_timeout	ofp_flow_entry_hard_timeout

#### 5.1.2 Enumerations

##### 5.1.2.1 ofp\_flow\_mod\_command

Table 5.2. enum ofp\_flow\_mod\_command

Name	Description
OFPPFC_ADD	New flow.
OFPPFC_MODIFY	Modify all matching flows.
OFPPFC_MODIFY_STRICT	Modify entry strictly matching wildcards and priority.
OFPPFC_DELETE	Delete all matching flows.
OFPPFC_DELETE_STRICT	Delete entry strictly matching wildcards and priority.
STATS_FAKE_CMD	Fake command to be used when creating a flow entry to be used by ofp_flow_entry_stats().

##### 5.1.2.2 ofp\_flow\_table\_type

Table 5.3. enum ofp\_flow\_table\_type

Name	Description
OFP_FLOW_TABLE_OVS_OPENFLOW	Open vSwitch table via OpenFlow protocol.
OFP_FLOW_TABLE_NFP_CACHE	NFP hardware flow cache (not implemented).

### 5.1.2.3 ofp\_group\_no

**Table 5.4. enum ofp\_group\_no**

Name	Description
OFPG_MAX	Last usable group number.
OFPG_ALL	Represents all groups for group delete commands.
OFPG_ANY	Wildcard group used only for flow stats requests. Selects all flows regardless of group (including flows with no group).

### 5.1.2.4 ofp\_port\_no

**Table 5.5. enum ofp\_port\_no**

Name	Description
OFPP_MAX	Maximum number of physical and logical switch ports.
OFPP_IN_PORT	Send the packet out the input port. This reserved port must be explicitly used in order to send back out of the input port.
OFPP_TABLE	Submit the packet to the first flow table NB: This destination port can only be used in packet-out messages.
OFPP_NORMAL	Process with normal L2/L3 switching.
OFPP_FLOOD	All physical ports in VLAN, except input port and those blocked or link down.
OFPP_ALL	All physical ports except input port.
OFPP_CONTROLLER	Send to controller.
OFPP_LOCAL	Local openflow "port".
OFPP_ANY	Wildcard port used only for flow mod (delete) and flow stats requests. Selects all flows regardless of output port (including flows with no output port).

### 5.1.2.5 ofp\_table

**Table 5.6. enum ofp\_table**

Name	Description
OFPTT_MAX	Last usable table number.
OFPTT_ALL	Wildcard table used for table config, flow stats and flow deletes.

## 5.1.3 Typedefs

### 5.1.3.1 openflow\_error\_fptr\_t

Error callback function pointer.

OpenFlow error indications are sent asynchronously by the switch. Registering this error callback in the `ofp_sdn_config` struct passed to `ofp_sdn_sys_init()` allows the user to be notified of these errors. The `user_context` (also initially configured in `ofp_sdn_config`) is given back to the user, along with `xid` (a transaction ID which can be set on any flow entry using the `ofp_flow_entry_xid()` function). The `errtype` and `errcode` are received from the switch, and refer to OpenFlow errors which can be found in the OpenFlow specification (Section 7.4.4 in OpenFlow 1.3.4). The type value indicates the high-level type of error, while the code value is interpreted based on the type.

**Table 5.7. typedef openflow\_error\_fptr\_t**

Type	Definition
<code>openflow_error_fptr_t</code>	<code>void(* openflow_error_fptr_t)(uint64_t user_context, uint32_t xid, uint16_t errt</code>

**See Also:**

- `ofp_sdn_barrier()`

## 5.1.4 Functions

### 5.1.4.1 ofp\_sdn\_sys\_init

**Prototype:**

```
struct ofp_sdn* ofp_sdn_sys_init(struct ofp_sdn_config* conf)
```

**Description:**

`ofp_sdn` API system initialization.

This function needs to be called before any flow modifications can be performed. The resulting handle can be passed repeatedly to `ofp_flow_entry_create()`, before eventually calling `ofp_sdn_sys_shutdown()` when done.

Open vSwitch needs to be configured to allow controller connections, e.g. `ovs-vsctl set-controller br0 tcp:6653`

**Table 5.8. ofp\_sdn\_sys\_init parameters**

Name	Description
<i>conf</i>	Pointer to struct containing config. If <code>ip_addr[0]==0</code> , the <code>/etc/netronome.conf</code> file will be parsed, and if that file is not found, IP 127.0.0.1 and port 6653 will be used.

**Returns:**



The initialized handle is returned on success, NULL otherwise.

**See Also:**

- ofp\_sdn\_sys\_shutdown()
- ofp\_sdn\_entry\_create()

### 5.1.4.2 ofp\_sdn\_sys\_shutdown

**Prototype:**

```
void ofp_sdn_sys_shutdown(struct ofp_sdn* nssdn)
```

**Description:**

API system cleanup

Reverts all the initialization done by ofp\_sdn\_sys\_init(), this must be called when done with the API.

### 5.1.4.3 ofp\_sdn\_barrier

**Prototype:**

```
int ofp_sdn_barrier(struct ofp_sdn* nssdn)
```

**Description:**

OpenFlow transaction barrier.

This call sends a barrier request and waits for a barrier reply. When the reply has been received, the user can be assured that all previous asynchronous flow requests have been completed.

**Table 5.9. ofp\_sdn\_barrier parameters**

Name	Description
<i>nssdn</i>	ofp_sdn API handle

**Returns:**

non-zero indicates failure

### 5.1.4.4 ofp\_sdn\_barrier\_timeout

**Prototype:**

```
int ofp_sdn_barrier_timeout(struct ofp_sdn* nssdn, uint32_t timeout_ms)
```

**Description:**

OpenFlow transaction barrier with timeout.

This call sends a barrier request and waits (with timeout) for a barrier reply. When the reply has been received, the user can be assured that all previous asynchronous flow requests have been completed.

**Table 5.10. ofp\_sdn\_barrier\_timeout parameters**

Name	Description
<i>nssdn</i>	ofp_sdn API handle
<i>timeout_ms</i>	Timeout period in milliseconds

**Returns:**

non-zero indicates failure or timeout

### 5.1.4.5 ofp\_flow\_entry\_create

**Prototype:**

```
struct ofp_flow_entry* ofp_flow_entry_create(struct ofp_sdn* nssdn, enum
ofp_flow_table_type ttype, enum ofp_flow_mod_command cmd, unsigned flags)
```

**Description:**

Create Flow entry handle

Allocates memory and initializes a Flow entry handle to be used with all match APIs. Initializes 'match' as a "catch-all" match that matches every packet. The returned handle must eventually be committed, or freed.



**Note**

The distinction between DELETE and DELETE\_STRICT, as well as between MODIFY and MODIFY\_STRICT, can be found in the OpenFlow specification as follows: Modify and delete flow mod commands have non-strict versions (OFPFC\_MODIFY and OFPFC\_DELETE) and strict versions (OFPFC\_MODIFY\_STRICT or OFPFC\_DELETE\_STRICT). In the strict versions, the set of match fields, all match fields, including their masks, and the priority, are strictly matched against the entry, and only an identical flow entry is modified or removed. For example, if a message to remove entries is sent that has no match fields included, the OFPFC\_DELETE command would delete all flow entries from the tables, while the OFPFC\_DELETE\_STRICT command would only delete a flow entry that applies to all packets at the specified priority

The meanings of the flag values can be found in the OpenFlow specification.

**Example (add a new flow entry to table 0 to match a specified IPv4 src address and send it out on the NORMAL virtual port):**

```
struct ofp_sdn *nssdn = ofp_sdn_sys_init(argv[0]); assert(nssdn);
struct ofp_flow_entry *fm =
    ofp_flow_entry_create(nssdn, OFP_FLOW_TABLE_OVS_OPENFLOW, OFPFC_ADD, 0); assert(fm);
```

```
ofp_flow_entry_match_field(fm, OFPXMT_OFB_IPV4_SRC, &val);
struct ofp_actions *act = ofp_action_create(OFP_ACTION_SET);
ofp_action_output(act, OFPP_NORMAL);
ofp_flow_entry_instr_write_actions(fm, act);
ofp_flow_entry_commit(fm);
ofp_sdn_sys_shutdown(nsflow);
```

**Example (delete all flow entries from table 0):**

```
struct ofp_flow_entry *fm =
    ofp_flow_entry_create(nsflow, OFP_FLOW_TABLE_OVS_OPENFLOW, OFPFC_DELETE, 0); assert(fm);
ofp_flow_entry_commit(fm, 0);
```

**Table 5.11. ofp\_flow\_entry\_create parameters**

Name	Description
<i>nsdn</i>	ofp_sdn API handle
<i>ttype</i>	enum ofp_flow_table_type (OFP_FLOW_TABLE_OVS_OPENFLOW, OFP_FLOW_TABLE_NFP_CACHE)
<i>cmd</i>	enum ofp_flow_mod_command (OFPFC_ADD, OFPFC_MODIFY, OFPFC_MODIFY_STRICT, OFPFC_DELETE, OFPFC_DELETE_STRICT)
<i>flags</i>	Combination of values from enum ofp_flow_mod_flags, ofp10_flow_mod_flags, ofp12_flow_mod_flags or ofp13_flow_mod_flags: OFPFF_SEND_FLOW_REM, OFPFF_CHECK_OVERLAP, OFPFF10_EMERG, OFPFF_RESET_COUNTS, OFPFF_NO_PKT_COUNTS, OFPFF_NO_BYT_COUNTS

**Returns:**

ofp\_flow\_entry handle

**See Also:**

- ofp\_sdn\_sys\_init()
- ofp\_flow\_entry\_commit()
- ofp\_flow\_entry\_free()

## 5.1.4.6 ofp\_flow\_entry\_free

**Prototype:**

```
void ofp_flow_entry_free(struct ofp_flow_entry* fe)
```

**Description:**

Free allocated memory for flow mod buffer. Used to destroy a flow handle which has not been committed.

**Table 5.12. ofp\_flow\_entry\_free parameters**

Name	Description
<i>fe</i>	Flow entry handle

**See Also:**

- `ofp_flow_entry_create()`

### 5.1.4.7 ofp\_flow\_entry\_match\_field

**Prototype:**

```
int ofp_flow_entry_match_field(struct ofp_flow_entry* fe, enum oxm_ofb_match_fields
match_field, const union action_value* value)
```

**Description:**

Add match field to flow mod (without wildcard mask).

**Table 5.13. ofp\_flow\_entry\_match\_field parameters**

Name	Description
<i>fe</i>	Flow entry handle
<i>match_field</i>	Member of enum <code>oxm_ofb_match_fields</code> specifying field to match
<i>value</i>	Pointer to value of match field. All of the 16 and 32bit network packet fields must be specified in network byte order. Other fields such as metadata, cookie, port numbers must be specified in x86 host byte order.

**Returns:**

non-zero indicates failure

**See Also:**

- `ofp_flow_entry_create()`

### 5.1.4.8 ofp\_flow\_entry\_match\_field\_mask

**Prototype:**

```
int ofp_flow_entry_match_field_mask(struct ofp_flow_entry* fe, enum
oxm_ofb_match_fields match_field, const union action_value* value, const union
action_value* mask_value)
```

**Description:**

Adds match field to flow mod (with wildcard mask).

**Table 5.14. ofp\_flow\_entry\_match\_field\_mask parameters**

Name	Description
<i>fe</i>	Flow entry handle
<i>match_field</i>	Member of enum <code>oxm_ofb_match_fields</code> specifying field to match
<i>value</i>	Pointer to value of match field. All of the 16 and 32bit network packet fields must be specified in network byte order. Other fields such as metadata, cookie, port numbers must be specified in x86 host byte order.
<i>mask_value</i>	Pointer to mask value for <code>match_field</code> value

**Returns:**

non-zero indicates failure

**See Also:**

- `ofp_flow_entry_create()`

### 5.1.4.9 ofp\_flow\_entry\_cookie

**Prototype:**

```
int ofp_flow_entry_cookie(struct ofp_flow_entry* fe, uint64_t cookie, uint64_t mask)
```

**Description:**

Populates Cookie field in flow entry



**Note**

Modify and delete commands can also be filtered by cookie value, if the `cookie_mask` field contains a value other than 0. This constraint is that the bits specified by the `cookie_mask` in both the cookie field of the flow mod and a flow entry's cookie value must be equal. In other words,  $(\text{flow\_entry.cookie} \& \text{flow mod.cookie\_mask}) == (\text{flow mod.cookie} \& \text{flow mod.cookie\_mask})$ .

**Table 5.15. ofp\_flow\_entry\_cookie parameters**

Name	Description
<i>fe</i>	Flow entry handle
<i>cookie</i>	Cookie value, in x86 host byte order
<i>mask</i>	Cookie mask, only relevant for modify and delete flow entries

**Returns:**

non-zero indicates failure

**See Also:**

- `ofp_flow_entry_create()`

#### 5.1.4.10 ofp\_flow\_entry\_xid

**Prototype:**

```
int ofp_flow_entry_xid(struct ofp_flow_entry* fe, uint32_t xid)
```

**Description:**

Set the transaction ID. This is an arbitrary number associated with a transaction. This number is included in any responses, such as error callback, to facilitate pairing. It is thus recommended to be set to a unique value per transaction. It defaults to 0 when a new `ofp_sdn` instance is created, and auto-increments with each transaction. This function resets the current value, which auto-increments again for subsequent transactions if set again.

**Table 5.16. ofp\_flow\_entry\_xid parameters**

Name	Description
<i>fe</i>	Flow entry handle
<i>xid</i>	Transaction ID

**Returns:**

non-zero indicates failure

**See Also:**

- `ofp_flow_entry_create()`

#### 5.1.4.11 ofp\_flow\_entry\_hard\_timeout

**Prototype:**

```
int ofp_flow_entry_hard_timeout(struct ofp_flow_entry* fe, uint16_t hard_timeout)
```

**Description:**

Populates Hard Timeout in flow entry (for temporary flow entries which can time out)

**Table 5.17. ofp\_flow\_entry\_hard\_timeout parameters**

Name	Description
<i>fe</i>	Flow entry handle
<i>hard_timeout</i>	Hard timeout Value

**Returns:**

non-zero indicates failure

**See Also:**

- `ofp_flow_entry_create()`

### 5.1.4.12 ofp\_flow\_entry\_idle\_timeout

**Prototype:**

```
int ofp_flow_entry_idle_timeout(struct ofp_flow_entry* fe, uint16_t idle_timeout)
```

**Description:**

Populates Idle Timeout in flow entry (for temporary flow entries which can time out)

**Table 5.18. ofp\_flow\_entry\_idle\_timeout parameters**

Name	Description
<i>fe</i>	Flow entry handle
<i>idle_timeout</i>	Idle timeout Value

**Returns:**

non-zero indicates failure

**See Also:**

- `ofp_flow_entry_create()`

### 5.1.4.13 ofp\_flow\_entry\_priority

**Prototype:**

```
int ofp_flow_entry_priority(struct ofp_flow_entry* fe, uint16_t priority)
```

**Description:**

Populates Priority in flow entry

**Table 5.19. ofp\_flow\_entry\_priority parameters**

Name	Description
<i>fe</i>	Flow entry handle
<i>priority</i>	Priority Value

**Returns:**

non-zero indicates failure

**See Also:**

- `ofp_flow_entry_create()`

#### 5.1.4.14 ofp\_flow\_entry\_out\_port

##### Prototype:

```
int ofp_flow_entry_out_port(struct ofp_flow_entry* fe, uint32_t out_port)
```

##### Description:

Populates `out_port` in flow entry

"The `out_port` and `out_group` fields optionally filter the scope of `OFPPFC_DELETE` and `OFPPFC_DELETE_STRICT` messages by output port and group. If either `out_port` or `out_group` contains a value other than `OFPP_ANY` or `OFPG_ANY` respectively, it introduces a constraint when matching. This constraint is that the flow entry must contain an output action directed at that port or group. Other constraints such as `ofp_match` structs and priorities are still used; this is purely an additional constraint. Note that to disable output filtering, both `out_port` and `out_group` must be set to `OFPP_ANY` and `OFPG_ANY` respectively. These fields are ignored by `OFPPFC_ADD`, `OFPPFC_MODIFY` or `OFPPFC_MODIFY_STRICT` messages."

**Table 5.20. ofp\_flow\_entry\_out\_port parameters**

Name	Description
<i>fe</i>	Flow entry handle
<i>out_port</i>	Output Port Number

##### Returns:

non-zero indicates failure

##### See Also:

- `ofp_flow_entry_create()`

#### 5.1.4.15 ofp\_flow\_entry\_out\_group

##### Prototype:

```
int ofp_flow_entry_out_group(struct ofp_flow_entry* fe, uint32_t out_group)
```

##### Description:

Populates `out_group` in flow entry

"The `out_port` and `out_group` fields optionally filter the scope of `OFPPFC_DELETE` and `OFPPFC_DELETE_STRICT` messages by output port and group. If either `out_port` or `out_group` contains a value other than `OFPP_ANY` or `OFPG_ANY` respectively, it introduces a constraint when matching. This constraint is that the flow entry must contain an output action directed at that port or group. Other constraints such as `ofp_match` structs and priorities are still used; this is purely an additional constraint. Note that to disable output



filtering, both `out_port` and `out_group` must be set to `OFPP_ANY` and `OFPG_ANY` respectively. These fields are ignored by `OFPPC_ADD`, `OFPPC_MODIFY` or `OFPPC_MODIFY_STRICT` messages."

**Table 5.21. `ofp_flow_entry_out_group` parameters**

Name	Description
<i>fe</i>	Flow entry handle
<i>out_group</i>	Output Group Number

**Returns:**

non-zero indicates failure

**See Also:**

- `ofp_flow_entry_create()`

### 5.1.4.16 `ofp_flow_entry_table_num`

**Prototype:**

```
int ofp_flow_entry_table_num(struct ofp_flow_entry* fe, uint8_t tablenum)
```

**Description:**

Populates table num in flow entry

**Table 5.22. `ofp_flow_entry_table_num` parameters**

Name	Description
<i>fe</i>	Flow entry handle
<i>tablenum</i>	Table number; <code>OFPTT_ALL</code> can be used when deleting flow entries to indicate that matching flow entries are to be deleted from all flow tables. Note that 254 is reserved for internal use by Open vSwitch.

**Returns:**

non-zero indicates failure

**See Also:**

- `ofp_flow_entry_create()`

### 5.1.4.17 `ofp_flow_entry_instr_goto_table`

**Prototype:**

```
int ofp_flow_entry_instr_goto_table(struct ofp_flow_entry* fe, uint8_t table_id)
```

### Description:

Adds Goto Table instruction to flow entry

**Table 5.23. ofp\_flow\_entry\_instr\_goto\_table parameters**

Name	Description
<i>fe</i>	Flow entry handle
<i>table_id</i>	Table ID value

### Returns:

non-zero indicates failure

### See Also:

- `ofp_flow_entry_create()`

## 5.1.4.18 ofp\_flow\_entry\_instr\_write\_metadata

### Prototype:

```
int ofp_flow_entry_instr_write_metadata(struct ofp_flow_entry* fe, uint64_t metadata,
uint64_t metadata_mask)
```

### Description:

Adds Write Metadata instruction to flow entry (with metadata\_mask value)

**Table 5.24. ofp\_flow\_entry\_instr\_write\_metadata parameters**

Name	Description
<i>fe</i>	Flow entry handle
<i>metadata</i>	Metadata value, in x86 host byte order
<i>metadata_mask</i>	Metadata mask value, ((uint64_t)(-1)) to match all bits

### Returns:

non-zero indicates failure

### See Also:

- `ofp_flow_entry_create()`

## 5.1.4.19 ofp\_flow\_entry\_instr\_clear\_actions

### Prototype:

```
int ofp_flow_entry_instr_clear_actions(struct ofp_flow_entry* fe)
```

**Description:**

Adds Clear Actions instruction to flow entry, to clear the Action Set associated with the packet.



**Note**

In OpenFlow semantics, the Clear-Actions instruction will always be executed prior to the Write-Actions instruction.

**Table 5.25. ofp\_flow\_entry\_instr\_clear\_actions parameters**

Name	Description
<i>fe</i>	Flow entry handle

**Returns:**

non-zero indicates failure

**See Also:**

- ofp\_flow\_entry\_create()

## 5.1.4.20 ofp\_flow\_entry\_instr\_meter

**Prototype:**

```
int ofp_flow_entry_instr_meter(struct ofp_flow_entry* fe, uint32_t meter_id)
```

**Description:**

Adds Meter instruction to flow entry

**Table 5.26. ofp\_flow\_entry\_instr\_meter parameters**

Name	Description
<i>fe</i>	Flow entry handle
<i>meter_id</i>	Meter ID value, in x86 host byte order

**Returns:**

non-zero indicates failure

**See Also:**

- ofp\_flow\_entry\_create()

## 5.1.4.21 ofp\_flow\_entry\_instr\_write\_actions

**Prototype:**

```
int ofp_flow_entry_instr_write_actions(struct ofp_flow_entry* fe, struct ofp_actions*
act)
```

### Description:

Add the Write-Actions instruction to the flow entry, which merges the specified action set into the current action set associated with the packet. The action set is specified by calls to `ofp_action_*` APIs while in the action set mode, the order in which those APIs are called is not significant (unless the same action is (incorrectly) specified more than once, in which case the last one will be the only one that takes effect).



### Note

In accordance with OpenFlow semantics, if a user wishes to completely specify the action set instead of merging, the Clear-Actions instruction should be enabled on the flow entry in addition to the Write-Actions instruction. The Clear-Actions instruction will always be executed before the Write-Actions instruction.

The action set accumulated on a packet after matching (possibly more than one) flow entries in flow tables is only executed when the final matching flow entry does not contain a Goto-Table instruction. The order of execution of these actions is determined by the OpenFlow specification. Contrast this with the Apply-Actions instruction, where an ordered action list is executed as soon as a packet matches a flow entry.

**Table 5.27. `ofp_flow_entry_instr_write_actions` parameters**

Name	Description
<i>fe</i>	Flow entry handle
<i>act</i>	Actions handle, ownership transfers to this API (caller must not access or free)

### Returns:

non-zero indicates failure

### See Also:

- `ofp_action_create()`

## 5.1.4.22 `ofp_flow_entry_instr_apply_actions`

### Prototype:

```
int ofp_flow_entry_instr_apply_actions(struct ofp_flow_entry* fe, struct ofp_actions*
act)
```

### Description:

Add the Apply-Actions instruction to the flow entry. The action list is populated by calls to `ofp_action_*` APIs while in the action list mode, and sequential order of those API calls is significant.



## Note

When a packet matches a flow entry with an Apply-Actions instruction, the ordered action list is executed immediately. Contrast this with the Write-Actions instruction, which only executes actions after the last matching flow entry does not have a Goto-Table instruction.

**Table 5.28. ofp\_flow\_entry\_instr\_apply\_actions parameters**

Name	Description
<i>fe</i>	Flow entry handle
<i>act</i>	Actions handle, ownership transfers to this API (caller must not access or free)

### Returns:

non-zero indicates failure

### See Also:

- ofp\_action\_create()

## 5.1.4.23 ofp\_flow\_entry\_commit

### Prototype:

```
int ofp_flow_entry_commit(struct ofp_flow_entry* fe)
```

### Description:

Commit the flow entry to the specified table number.

**Table 5.29. ofp\_flow\_entry\_commit parameters**

Name	Description
<i>fe</i>	Flow entry handle

### Returns:

non-zero indicates failure

### See Also:

- ofp\_flow\_entry\_create()

## 5.1.4.24 ofp\_flow\_entry\_stats

### Prototype:

```
int ofp_flow_entry_stats(struct ofp_flow_entry* fe, uint64_t* bytes, uint64_t* packets,
uint32_t* flows)
```

### Description:

Request flow statistics. The total bytes, packets and number of flows matching the criteria in *fe* are returned. By default *table\_id* is OFPPTT\_ALL, *out\_port* is OFPP\_ANY and *out\_group* is OFPG\_ANY. Only the match fields in *fe* are taken into consideration, any instructions+actions that may have been specified are ignored.

**Table 5.30. ofp\_flow\_entry\_stats parameters**

Name	Description
<i>fe</i>	Flow entry handle, must be created with the special STATS_FAKE_CMD command type
<i>bytes</i>	Pointer to variable in which resulting bytes are stored
<i>packets</i>	Pointer to variable in which resulting packets are stored
<i>flows</i>	Pointer to variable in which resulting flows are stored

### Returns:

non-zero indicates failure

## 5.2 ofp\_group

### 5.2.1 Enumerations

#### 5.2.1.1 ofp\_group\_mod\_command

**Table 5.31. enum ofp\_group\_mod\_command**

Name	Description
OFPGC_ADD	
OFPGC_MODIFY	
OFPGC_DELETE	

#### 5.2.1.2 ofp\_group\_type

**Table 5.32. enum ofp\_group\_type**

Name	Description
OFPGT_ALL	
OFPGT_SELECT	
OFPGT_INDIRECT	

Name	Description
OFPGT_FF	

## 5.2.2 Functions

### 5.2.2.1 ofp\_group\_create

**Prototype:**

```
struct ofp_group* ofp_group_create(struct ofp_sdn* nssdn, enum ofp_group_mod_command cmd, enum ofp_group_type typ, uint32_t id)
```

**Description:**

Create a new group modifier. The returned handle must eventually be committed, or freed.

**Table 5.33. ofp\_group\_create parameters**

Name	Description
<i>nssdn</i>	
<i>cmd</i>	Add, modify or delete.
<i>typ</i>	Group type.
<i>id</i>	The group identifier.

**Returns:**

struct ofp\_group\* handle to create group modifier

### 5.2.2.2 ofp\_group\_add\_action\_bucket

**Prototype:**

```
int ofp_group_add_action_bucket(struct ofp_group* grp, uint16_t weight, uint32_t watch_port, uint32_t watch_group, struct ofp_actions* actions)
```

**Description:**

Add a new bucket to the list, with specified weight and action set.

**Table 5.34. ofp\_group\_add\_action\_bucket parameters**

Name	Description
<i>grp</i>	Group modifier handle.

Name	Description
<i>weight</i>	Relative weight of bucket.
<i>watch_port</i>	Port whose state affects this bucket (only for fast failover groups).
<i>watch_group</i>	Group whose state affects this bucket (only for fast failover groups).
<i>actions</i>	Action set created by ofp_action APIs.

**Returns:**

int Non-zero on failure.

### 5.2.2.3 ofp\_group\_commit

**Prototype:**

```
int ofp_group_commit(struct ofp_group* grp)
```

**Description:**

Commit the group modifier.

**Table 5.35. ofp\_group\_commit parameters**

Name	Description
<i>grp</i>	Group modifier handle.

**Returns:**

int Non-zero on failure.

### 5.2.2.4 ofp\_group\_free

**Prototype:**

```
int ofp_group_free(struct ofp_group* grp)
```

**Description:**

Free the group modifier. Used only on group handles that haven't been committed.

**Table 5.36. ofp\_group\_free parameters**

Name	Description
<i>grp</i>	Group modifier handle.

**Returns:**

int Non-zero on failure.



## 5.3 ofp\_action

### 5.3.1 Defines

**Table 5.37. ofp\_action.h Defines**

Defined	Definition
ETH_ADDR_LEN	6

### 5.3.2 Enumerations

#### 5.3.2.1 ofp\_action\_mode

**Table 5.38. enum ofp\_action\_mode**

Name	Description
OFP_ACTION_SET	Add subsequent actions to an Action Set, as used by the Write-Actions instruction.
OFP_ACTION_LIST	Add subsequent actions to an Action List, as used by the Apply-Actions instruction.

#### 5.3.2.2 oxm\_ofb\_match\_fields

**Table 5.39. enum oxm\_ofb\_match\_fields**

Name	Description
OFPXMT_OFB_IN_PORT	
OFPXMT_OFB_IN_PHY_PORT	
OFPXMT_OFB_METADATA	
OFPXMT_OFB_ETH_DST	
OFPXMT_OFB_ETH_SRC	
OFPXMT_OFB_ETH_TYPE	
OFPXMT_OFB_VLAN_VID	
OFPXMT_OFB_VLAN_PCP	
OFPXMT_OFB_IP_DSCP	
OFPXMT_OFB_IP_ECN	
OFPXMT_OFB_IP_PROTO	

Name	Description
OFPXMT_OFB_IPV4_SRC	
OFPXMT_OFB_IPV4_DST	
OFPXMT_OFB_TCP_SRC	
OFPXMT_OFB_TCP_DST	
OFPXMT_OFB_UDP_SRC	
OFPXMT_OFB_UDP_DST	
OFPXMT_OFB_SCTP_SRC	
OFPXMT_OFB_SCTP_DST	
OFPXMT_OFB_ICMPV4_TYPE	
OFPXMT_OFB_ICMPV4_CODE	
OFPXMT_OFB_ARP_OP	
OFPXMT_OFB_ARP_SPA	
OFPXMT_OFB_ARP_TPA	
OFPXMT_OFB_ARP_SHA	
OFPXMT_OFB_ARP_THA	
OFPXMT_OFB_IPV6_SRC	
OFPXMT_OFB_IPV6_DST	
OFPXMT_OFB_IPV6_FLABEL	
OFPXMT_OFB_ICMPV6_TYPE	
OFPXMT_OFB_ICMPV6_CODE	
OFPXMT_OFB_IPV6_ND_TARGET	
OFPXMT_OFB_IPV6_ND_SLL	
OFPXMT_OFB_IPV6_ND_TLL	
OFPXMT_OFB_MPLS_LABEL	
OFPXMT_OFB_MPLS_TC	
OFPXMT_OFB_MPLS_BOS	
OFPXMT_OFB_PBB_ISID	
OFPXMT_OFB_TUNNEL_ID	
OFPXMT_OFB_IPV6_EXTHDR	

## 5.3.3 Functions

### 5.3.3.1 ofp\_action\_create

**Prototype:**

```
struct ofp_actions* ofp_action_create(enum ofp_action_mode a)
```

#### Description:

Create an action handle for use with `ofp_flow_entry` or `ofp_group` APIs. The supplied mode is used to distinguish between adding actions to a set or an ordered list.

**Table 5.40. ofp\_action\_create parameters**

Name	Description
<i>a</i>	ofp_action_mode enum

#### See Also:

- `ofp_flow_entry_instr_write_actions()`
- `ofp_flow_entry_instr_apply_actions()`
- `ofp_group_mod_add_action_bucket()`

### 5.3.3.2 ofp\_action\_free

#### Prototype:

```
void ofp_action_free(struct ofp_actions* act)
```

#### Description:

Free allocated memory for actions buffer. Only necessary if the action handle has not had ownership transferred to `ofp_flow_entry` or `ofp_group` API instances.

**Table 5.41. ofp\_action\_free parameters**

Name	Description
<i>act</i>	Action handle

#### See Also:

- `ofp_action_create()`

### 5.3.3.3 ofp\_action\_output

#### Prototype:

```
int ofp_action_output(struct ofp_actions* act, uint32_t output_port)
```

#### Description:

Adds Output action. The output port value can also include the special OpenFlow values: `OFPP_IN_PORT` Send the packet back out the input port. `OFPP_NORMAL` Permit normal forwarding to select the output port. `OFPP_CONTROLLER` Send to controller.

**Table 5.42. ofp\_action\_output parameters**

Name	Description
<i>act</i>	Action handle
<i>output_port</i>	Output port value in host byte order (x86 little endian)

**Returns:**

non-zero indicates failure

**See Also:**

- ofp\_action\_create()

### 5.3.3.4 ofp\_action\_group

**Prototype:**

```
int ofp_action_group(struct ofp_actions* act, uint32_t group_id)
```

**Description:**

Adds Group action.

**Table 5.43. ofp\_action\_group parameters**

Name	Description
<i>act</i>	Action handle
<i>group_id</i>	Group ID value in host byte order (x86 little endian)

**Returns:**

non-zero indicates failure

**See Also:**

- ofp\_action\_create(), ofp\_group\_create()

### 5.3.3.5 ofp\_action\_mpls\_ttl

**Prototype:**

```
int ofp_action_mpls_ttl(struct ofp_actions* act, uint8_t mpls_ttl)
```

**Description:**

Add Set MPLS TTL action

**Table 5.44. ofp\_action\_mpls\_ttl parameters**

Name	Description
<i>act</i>	Action handle
<i>mpls_ttl</i>	mpls_ttl value

**Returns:**

non-zero indicates failure

**See Also:**

- ofp\_action\_create()

### 5.3.3.6 ofp\_action\_dec\_mpls\_ttl

**Prototype:**

```
int ofp_action_dec_mpls_ttl(struct ofp_actions* act)
```

**Description:**

Add Decrement MPLS TTL action

**Table 5.45. ofp\_action\_dec\_mpls\_ttl parameters**

Name	Description
<i>act</i>	Action handle

**Returns:**

non-zero indicates failure

**See Also:**

- ofp\_action\_create()

### 5.3.3.7 ofp\_action\_ip\_ttl

**Prototype:**

```
int ofp_action_ip_ttl(struct ofp_actions* act, uint8_t ip_ttl)
```

**Description:**

Add Set IP TTL action

**Table 5.46. ofp\_action\_ip\_ttl parameters**

Name	Description
<i>act</i>	Action handle

Name	Description
<i>ip_ttl</i>	ip_ttl value

**Returns:**

non-zero indicates failure

**See Also:**

- ofp\_action\_create()

### 5.3.3.8 ofp\_action\_dec\_ip\_ttl

**Prototype:**

```
int ofp_action_dec_ip_ttl(struct ofp_actions* act)
```

**Description:**

Add Decrement IP TTL action

**Table 5.47. ofp\_action\_dec\_ip\_ttl parameters**

Name	Description
<i>act</i>	Action handle

**Returns:**

non-zero indicates failure

**See Also:**

- ofp\_action\_create()

### 5.3.3.9 ofp\_action\_push\_mpls

**Prototype:**

```
int ofp_action_push_mpls(struct ofp_actions* act, uint16_t ethertype)
```

**Description:**

Add Push MPLS action

**Table 5.48. ofp\_action\_push\_mpls parameters**

Name	Description
<i>act</i>	Action handle

Name	Description
<i>ethertype</i>	ethertype value, in network byte order

**Returns:**

non-zero indicates failure

**See Also:**

- ofp\_action\_create()

### 5.3.3.10 ofp\_action\_pop\_mpls

**Prototype:**

```
int ofp_action_pop_mpls(struct ofp_actions* act)
```

**Description:**

Add Pop MPLS action

**Table 5.49. ofp\_action\_pop\_mpls parameters**

Name	Description
<i>act</i>	Action handle

**Returns:**

non-zero indicates failure

**See Also:**

- ofp\_action\_create()

### 5.3.3.11 ofp\_action\_set\_queue

**Prototype:**

```
int ofp_action_set_queue(struct ofp_actions* act, uint32_t queue_id)
```

**Description:**

Add Set Queue action

**Table 5.50. ofp\_action\_set\_queue parameters**

Name	Description
<i>act</i>	Action handle

Name	Description
<i>queue_id</i>	queue_id value

**Returns:**

non-zero indicates failure

**See Also:**

- ofp\_action\_create()

### 5.3.3.12 ofp\_action\_push\_vlan

**Prototype:**

```
int ofp_action_push_vlan(struct ofp_actions* act, uint16_t ethertype)
```

**Description:**

Adds Push VLAN action

**Table 5.51. ofp\_action\_push\_vlan parameters**

Name	Description
<i>act</i>	Action handle
<i>ethertype</i>	ethertype value in network byte order (only htons(0x8100) is currently supported)

**Returns:**

non-zero indicates failure

**See Also:**

- ofp\_action\_create()

### 5.3.3.13 ofp\_action\_pop\_vlan

**Prototype:**

```
int ofp_action_pop_vlan(struct ofp_actions* act)
```

**Description:**

Adds Pop VLAN action

**Table 5.52. ofp\_action\_pop\_vlan parameters**

Name	Description
<i>act</i>	Action handle



### Returns:

non-zero indicates failure

### See Also:

- ofp\_action\_create()

## 5.3.3.14 ofp\_action\_set\_field

### Prototype:

```
int ofp_action_set_field(struct ofp_actions* act, enum oxm_ofb_match_fields field,
union action_value* value)
```

### Description:

Add a set field action

**Table 5.53. ofp\_action\_set\_field parameters**

Name	Description
<i>act</i>	Action handle
<i>field</i>	Member of enum oxm12_ofb_match_fields specifying field to set
<i>value</i>	Pointer to value of match field. All of the 16 and 32bit network packet fields must be specified in network byte order. Other fields such as metadata, cookie, port numbers must be specified in x86 host byte order.

### Returns:

non-zero indicates failure

## 5.4 ns\_sdn\_health

### 5.4.1 Defines

**Table 5.54. ns\_sdn\_health.h Defines**

Defined	Definition
NS_SDN_HEALTH_FLAG_ENABLE_COLOR_OUTPUT	Initialization flag: If specified, all text report generated by the library will be color coded for display purposes.
NS_SDN_HEALTH_FLAG_ENABLE_SYSLOG	1

Defined	Definition
	Initialization flag: If specified, all failures will be logged via syslog when 'ns_sdn_health_check' is called. Syslog messages are generated with a syslog id of "ns_sdn_health" and logged to the LOCAL7 facility.
NS_SDN_HEALTH_FLAG_FAIL_ON_WARN	4  Initialization flag: If specified, warnings are treated as failures.
NS_SDN_HEALTH_OPTIONAL_DISABLE_ALL	( 0 )
NS_SDN_HEALTH_OPTIONAL_ENABLE_ALL	( ~(UINT32_MAX << NS_SDN_HEALTH_OPTIONAL_MAX) )
NS_SDN_HEALTH_PID_ABS_FILE	NS_SDN_HEALTH_PID_PATH" / "NS_SDN_HEALTH_PID_FILE
NS_SDN_HEALTH_PID_FILE	"lock.pid"
NS_SDN_HEALTH_PID_PATH	"/tmp/ovs.lock.v2.2"

## 5.4.2 Enumerations

### 5.4.2.1 ns\_sdn\_health\_failure\_code\_t

**Table 5.55.** enum ns\_sdn\_health\_failure\_code\_t

Name	Description
NS_SDN_HEALTH_NO_FAILURE	No failure.
NS_SDN_HEALTH_OVSDB_SERVER_FAILED	OVSDb server process failed.
NS_SDN_HEALTH_OVS_VSWITCHD_FAILED	OVS VSWITCHD process failed.
NS_SDN_HEALTH_VIRTIORELAYD_FAILED	Virtiorelayd process failed.
NS_SDN_HEALTH_NFP_MODULE_NOT_LOADED	Missing NFP kernel module.
NS_SDN_HEALTH_CMSG_MODULE_NOT_LOADED	Missing NFP control message module.
NS_SDN_HEALTH_FALLBACK_MODULE_NOT_LOADED	Missing NFP Fallback module.
NS_SDN_HEALTH_OFFLOADS_MODULE_NOT_LOADED	Missing NFP Offload module.
NS_SDN_HEALTH_OVS_MODULE_NOT_LOADED	Missing Open vSwitch kernel module.
NS_SDN_HEALTH_FIRMWARE_NOT_LOADED	Firmware not loaded.
NS_SDN_HEALTH_NFP_FPC_UNRESPONSIVE	Flow processing cores are unresponsive.
NS_SDN_HEALTH_NFP_CTRL_CHAN_UNRESPONSIVE	NFP Control channel is unresponsive.
NS_SDN_HEALTH_TRAFFIC_DIVERTED	Traffic diverted to the fallback channel.
NS_SDN_HEALTH_NFP_INGRESS_NBI_BACKLOGGED	Ingress NBI processing cores are possibly locked up.
NS_SDN_HEALTH_NFP_NOT_DETECTED	NFP not detected.
NS_SDN_HEALTH_ECC_ERRORS_DETECTED	ECC hardware errors detected.
NS_SDN_HEALTH_PARITY_ERRORS_DETECTED	Parity hardware errors detected.

Name	Description
NS_SDN_HEALTH_CONFIGURATOR_CHECK	Check for up to date configurator.
NS_SDN_HEALTH_ERR47_WORKAROUND	ERR47 workaround not detected.
NS_SDN_HEALTH_MAX	Marker to indicate maximum number of enums.

### 5.4.2.2 ns\_sdn\_health\_optional\_checks\_t

**Table 5.56.** enum ns\_sdn\_health\_optional\_checks\_t

Name	Description
NS_SDN_HEALTH_OPTIONAL_VIRTIORELAYD	Virtiorelayd process check.
NS_SDN_HEALTH_OPTIONAL_MAX	Marker to indicate maximum number of enums.

## 5.4.3 Typedefs

### 5.4.3.1 ns\_sdn\_health\_failure\_t

**Table 5.57.** typedef ns\_sdn\_health\_failure\_t

Type	Definition
ns_sdn_health_failure_t	struct ns_sdn_health_failure*

## 5.4.4 Functions

### 5.4.4.1 ns\_sdn\_health\_init

**Prototype:**

```
ns_sdn_health_handle_t ns_sdn_health_init(unsigned int flags, unsigned int
enabled_optional_checks)
```

**Description:**

Initialize a new instance of the SDN health check API.

**Table 5.58.** ns\_sdn\_health\_init parameters

Name	Description
<i>flags</i>	Bitmap of initialization flags

Name	Description
<i>enabled_optional_checks</i>	Bitmap of optional checks which are enabled. Bitmap must be created with a '1' value bit shifted with the corresponding enum value in the 'ns_sdn_health_optional_checks_t' enum.

**Returns:**

New instance. Returns NULL if function call fails. In case of failure, sets ERRNO value appropriately. ERRNO value of 'EBUSY' indicates the existence of the PID file /tmp/sdn\_health.pid. The memory allocated needs to be freed with a call to 'ns\_sdn\_health\_destroy' once done.

## 5.4.4.2 ns\_sdn\_health\_destroy

**Prototype:**

```
void ns_sdn_health_destroy(ns_sdn_health_handle_t handle)
```

**Description:**

Destroys an instance of the SDN health check API.

**Table 5.59. ns\_sdn\_health\_destroy parameters**

Name	Description
<i>handle</i>	Health check handle

## 5.4.4.3 ns\_sdn\_health\_check

**Prototype:**

```
int ns_sdn_health_check(ns_sdn_health_handle_t handle)
```

**Description:**

Executes the health check on the SDN 2.2 system. This function should be executed periodically or whenever the current system health needs to be determined.

Previous health failures will be cleared with each call to this function.

**Table 5.60. ns\_sdn\_health\_check parameters**

Name	Description
<i>handle</i>	Health check handle

**Returns:**

The number of failures detected. Returns -1 when unable to perform health check due to errors. Warnings are not counted unless the 'NS\_SDN\_HEALTH\_FLAG\_FAIL\_ON\_WARN' flag is specified.

#### 5.4.4.4 ns\_sdn\_health\_failure\_count

**Prototype:**

```
int ns_sdn_health_failure_count(ns_sdn_health_handle_t handle)
```

**Description:**

Returns the number of failures detected during the last execution of 'ns\_sdn\_health\_check'.

**Table 5.61. ns\_sdn\_health\_failure\_count parameters**

Name	Description
<i>handle</i>	Health check handle

**Returns:**

The number of failures detected. Will return 0 if 'ns\_sdn\_health\_check' has not been executed yet. Warnings are not counted unless the 'NS\_SDN\_HEALTH\_FLAG\_FAIL\_ON\_WARN' flag is specified. Returns -1 when unable to perform health check due to errors.

#### 5.4.4.5 ns\_sdn\_health\_warning\_count

**Prototype:**

```
int ns_sdn_health_warning_count(ns_sdn_health_handle_t handle)
```

**Description:**

Returns the number of warnings detected during the last execution of 'ns\_sdn\_health\_check'.

**Table 5.62. ns\_sdn\_health\_warning\_count parameters**

Name	Description
<i>handle</i>	Health check handle

**Returns:**

The number of warnings detected. Will return 0 if 'ns\_sdn\_health\_check' has not been executed yet. Returns -1 when unable to perform health check due to errors.

#### 5.4.4.6 ns\_sdn\_health\_next\_failure

**Prototype:**

```
int ns_sdn_health_next_failure(ns_sdn_health_handle_t handle, ns_sdn_health_failure_t* fail, ns_sdn_health_failure_code_t* failcode, const char** failstr, const char** affected_component, const char** recommended_action)
```

### Description:

Returns the failure structure of the next available health check failure. The number of health check failures available is determined from the call to 'ns\_sdn\_health\_check'.

**Table 5.63. ns\_sdn\_health\_next\_failure parameters**

Name	Description
<i>handle</i>	Health check handle
<i>fail</i>	Output pointer to failure structure. Can be NULL if the caller doesn't want to receive this data. Caller must not free this pointer. Library sets this to NULL if there are no more failures available.
<i>failcode</i>	Output pointer to the failure code. Can be NULL if the caller doesn't want to receive this data.
<i>failstr</i>	Output pointer to failure string. Can be NULL if the caller doesn't want to receive this data. Caller must not free this pointer. Library sets this to NULL if there are no more failures available.
<i>affected_component</i>	Output pointer to string indicating affected component. Can be NULL if the caller doesn't want to receive this data. Caller must not free this pointer. Library sets this to NULL if there are no more failures available.
<i>recommended_action</i>	Output pointer to string indicating recommended action. Can be NULL if the caller doesn't want to receive this data. Caller must not free this pointer. Library sets this to NULL if there are no more failures available.

### Returns:

0 on success, 1 on no more failures and -1 on error

## 5.4.4.7 ns\_sdn\_health\_generate\_report

### Prototype:

```
int ns_sdn_health_generate_report(ns_sdn_health_handle_t handle, char* buffer, size_t
buffer_size, int verbose)
```

### Description:

Generate a health report.

**Table 5.64. ns\_sdn\_health\_generate\_report parameters**

Name	Description
<i>handle</i>	Health check handle
<i>buffer</i>	User provided buffer where report will be written to.
<i>buffer_size</i>	Size of the user provided buffer. Library will not write more than this many bytes to the buffer.

Name	Description
<i>verbose</i>	If non-zero writes details of which subcomponents failed if a health check failed, e.g. register names and values are written to the buffer.

**Returns:**

Returns positive value with amount of bytes written to buffer or -1 on error.

## 6. Technical Support

To obtain additional information, or to provide feedback, please email <[support@netronome.com](mailto:support@netronome.com)> or contact the nearest **Netronome** technical support representative.

### 6.1 Related Documents

Descriptive Name	Description
Netronome Network Flow Processor: Agilio OvS 2.2 Getting Started Guide	A guide to new users of Netronome's Agilio OvS software.
Netronome Network Flow Processor: Agilio OvS 2.2 Programmer's Reference Manual	A reference for programming and system design using the Agilio OvS software.
Intelligent Server Adapters: Hardware User Manual	Contains summary information on the Netronome Intelligent Server Adapter (ISA) PCIe card including card physical descriptions.
Netronome Network Flow Processor: Datasheet	Contains summary information on the Netronome Network Flow Processor NFP including a functional description, signal descriptions, electrical specifications, and mechanical specifications.
Netronome Network Flow Processor: Databook	Contains detailed reference information on the Netronome Network Flow Processor NFP.
Netronome Network Flow Processor: Development Tools User's Guide	Describes Programmer Studio and the development tools that can be accessed through Programmer Studio.
Netronome Network Flow Processor: Network Flow Assembler System User's Guide	Describes the syntax of the NFP's assembly language, supplies assembler usage information, and lists assembler warnings and errors.
Netronome Network Flow Processor: Microengine Programmer's Reference Manual	A reference for microcode programming of the Netronome Network Flow Processor NFP.
Netronome Network Flow Processor: Network Flow C Compiler User's Guide	Presents information, language structures and extensions to the language specific to the Netronome Network Flow C Compiler for Netronome NFP.
Netronome Network Flow Compiler LibC: Reference Manual	Specifies the subset and the extensions to the language that support the unique features of the Netronome Network Flow Processor NFP product line.
Open vSwitch Software Documentation	Agilio OvS software offers acceleration of Open vSwitch software. Refer to <a href="http://openvswitch.org/">http://openvswitch.org/</a> for more details on Open vSwitch.
OpenFlow Specification	Open vSwitch (which is accelerated by Agilio OvS software) is an OpenFlow switch implementation. Refer to <a href="https://www.opennetworking.org/sdn-">https://www.opennetworking.org/sdn-</a>



Descriptive Name	Description
	resources/openflow for more details on this specification.
Data Plane Development Kit Documentation	DPDK related documentation is available at <a href="http://dpdk.org">http://dpdk.org</a> .